

1. COMPANY OVERVIEW

Company Name: ThinkByte.AI
Website: <https://thinkbyte.ai/>
Founded: 2024
Headquarters: No. 235, 13th Cross, Hoysala Nagar 2nd Stage, Bangalore North, Bangalore, Karnataka, India - 560038
Company Type: Private, AI Consulting and Development
Funding Status: Unfunded / Bootstrapped
Team Size: Small managed team (30+ agent orchestration experts)
Google Analytics ID: G-RXK5XW514W

What ThinkByte Does

ThinkByte.AI is a managed AI development team specializing in LangGraph and multi-agent AI systems. They build production-grade AI agent solutions, providing agency-level reliability at startup speed. The company has delivered 20+ production AI systems across healthcare, HR tech, marketing/SEO, and construction/architecture verticals.

One-Liner: "We're a small team that specializes in LangGraph and AI agent development."

Elevator Pitch

"ThinkByte is a managed AI development team that builds LangGraph and AI agent solutions. Unlike solo freelancers who disappear, we provide agency-level reliability at startup speed. We've built 20+ production AI systems with measurable results: 90% automation rates, 60% cost reductions, and patent-pending innovations."

2. FOUNDERS AND LEADERSHIP

SriHarsha Thota - Co-Founder

Full Name: SriHarsha Thota
Role: Co-Founder, ThinkByte.AI
Location: Bengaluru, India
Education: B.E. in Civil Engineering, BITS Pilani (2009-2013)
Total Experience: 12+ years in technology and business leadership
LinkedIn: [linkedin.com/in/sriharsha-thota-37033616](https://www.linkedin.com/in/sriharsha-thota-37033616)
Angel Investor: Yes
Areas of Interest: ConsumerTech, Health, Travel, Mobility, Energy, Education, Automation, AI

Career Timeline:

- ThinkByte.AI - Co-Founder (January 2024 to Present): Leading AI consulting startup focused on LangGraph and multi-agent systems
- MediBuddy - Vice President, Consumer Growth (April 2022 to July 2024): Led consumer growth strategy for India's leading health tech platform
- MediBuddy - Consumer Business Head (March 2021 to March 2022): Delivered 15x business growth in a single year
- MediBuddy/DocsApp - Growth and Special Projects, Founder's Office (September 2019 to March 2021): Led business transformation and organizational restructuring post-merger
- DocsApp - Growth and Special Projects: Part of leadership team at telehealth platform before MediBuddy acquisition

- Earlier roles at Reforge, eGain Corporation, and Kalyani Global Engineering Pvt. Ltd.

Key Skills: Marketing, Leadership, Healthcare, JavaScript, SQL, Product Management, Business Scaling

Notable Quotes from SriHarsha Thota:

- "AI should empower businesses, not disrupt the customer journey."
- "We believe in augmenting human intelligence rather than replacing it."
- "We continuously track outcomes and adjust strategies proactively, even before pain points are experienced."

Vinay Kumar - Co-Founder

Full Name: Vinay Kumar

Role: Co-Founder, ThinkByte.AI

Education: BITS Pilani

LinkedIn: [linkedin.com/in/vinay-kumar-2b42a845](https://www.linkedin.com/in/vinay-kumar-2b42a845)

Focus: AI product development

Background: Vinay Kumar studied at BITS Pilani alongside SriHarsha Thota. He has experience shipping 20+ AI products that work in production. He serves as the hands-on technical co-founder driving engineering delivery, with a background in AI/ML engineering and product development.

3. MISSION, VISION AND PHILOSOPHY

Mission

To create collaborative, customer-centric strategies that not only drive success but also foster innovation and inclusivity. To help businesses effectively integrate AI solutions to streamline operations, improve decision-making, and identify new growth opportunities.

Vision

Position ThinkByte as a leader in ethical, practical, and scalable AI solutions that transform customer experiences by eliminating friction points and enabling seamless, personalized experiences.

Core Philosophy

1. Human-First AI

- AI should empower businesses, not disrupt the customer journey
- Focus on augmenting human intelligence rather than replacing it
- Human-in-loop quality control at critical decision points

2. Practical Innovation

- Solutions must deliver tangible value without disrupting existing processes
- Reject solutions that don't provide better outcomes than human-driven approaches
- First-principles problem-solving adaptable across industries

3. Ethical AI

- Transparency, fairness, accountability guide all operations
- Strict data privacy protocols and data anonymization
- Detailed audit trails for compliance

4. Long-Term Value

- Focus on sustained business outcomes over fleeting trends
- Continuous monitoring and proactive intervention
- Balance growth with customer satisfaction

5. Production-First Mindset

- Production-grade systems, not just POCs or demos
- Infrastructure, monitoring, and maintenance included
- 20+ production deployments as proof of capability

4. SERVICES AND ENGAGEMENT MODELS

Core Service Lines

Service 1: Production AI Agent Development

- Complete multi-agent systems from architecture to deployment
- 3-6 month timeline to production-ready platform
- Includes monitoring, error handling, cost optimization
- 20+ agent orchestration with LangGraph supervisor patterns

Service 2: AI Integration and Automation

- Embed AI agents into Microsoft Teams, Slack, CRM, ATS
- End-to-end workflow automation (80-90% coverage)
- Real-time event-driven architectures
- Bidirectional sync with existing enterprise systems

Service 3: Healthcare AI Innovation

- Patent-worthy innovation with clinical validation
- HIPAA-compliant architecture and data security
- Elderly/patient-accessible UX design
- Behavioral prediction models and adaptive content

Service 4: Enterprise AI Consulting

- Framework selection guidance (LangGraph vs AutoGen vs CrewAI)
- Architecture design for scalability and cost optimization
- Technical debt reduction and AI system rescue
- Migration from POC to production

Flexible Engagement Options

- Dedicated Team: Full managed team assigned to your project on monthly basis
- Fixed-Price MVP/Prototype: Scoped deliverables with milestone-based payments
- Code Review and Optimization: Review and improve existing AI systems
- Ongoing Support and Enhancement: Post-launch maintenance, monitoring, and feature additions

Engagement Model 1: Monthly FTE (Managed Teams)

- Pay monthly in advance, dedicated team members
- 160 hours/month per engineer (standard month)
- 9-hour availability window (client chooses slot within 09:00-21:00 IST)
- Minimum commitment: 6 months
- Payment: Invoice on 25th, due Net 7

- Scale down: 15 days notice
- Replacement SLA: 7 business days or less

Engagement Model 2: Hourly Retainer

- Prepaid retainer model, fund first week's hours upfront
- Minimum: 10 hours
- Weekly timesheets (Friday), invoice Wednesday, due in 2 days
- Auto-replenishment: Keep balance for upcoming week
- Work pauses if retainer depletes

Engagement Model 3: Project-Based (Fixed Scope)

- Scoped deliverables with milestones
- Minimum: \$5K
- Payment: 30-40% upfront, rest on milestones
- Change requests quoted and approved separately

What Is Included (Standard)

- Dedicated engineer time
- PM coordination
- Code reviews
- Basic DevOps
- Weekly progress reports
- Slack/Teams communication
- Git repository
- Basic documentation

What Is NOT Included

- Third-party API costs
- Cloud hosting (unless specified)
- LLM API costs (OpenAI, Anthropic)
- On-premises deployment
- 24/7 on-call support
- Overtime/weekend work (custom pricing available)
- Data preparation/labeling
- Training client's team (unless scoped)

5. TECHNOLOGY STACK

Agent Frameworks

- LangGraph: Primary specialization - "It's all we do"
- AutoGen: Full production experience
- Agents SDK: Active development
- LlamaIndex: RAG and indexing
- CrewAI: Multi-agent orchestration

LLM Integration

- OpenAI GPT-4: Primary LLM for agent reasoning
- Anthropic Claude: Alternative LLM, content generation
- Google Gemini: Multi-modal tasks
- Llama 3: Open-source deployments
- Custom Fine-Tuning: Domain-specific models

RAG and Vector Databases

- Pinecone: Managed vector search at scale
- Weaviate: Hybrid search, self-hosted
- Chroma: Lightweight local embeddings
- Qdrant: High-performance vector DB

Orchestration

- LangChain for chains and tool-calling
- LlamaIndex for indexing and retrieval
- Agent memory systems (persistent, short-term, long-term)
- Tool-calling APIs with structured outputs

Backend

- FastAPI (primary API framework)
- Node.js (real-time services)
- PostgreSQL (primary relational database)
- Supabase (backend-as-a-service)
- Redis (caching, queues, pub/sub)
- Celery (async task processing)
- Docker (containerization)
- Kubernetes (orchestration at scale)
- AWS (EC2, Lambda, S3)

Frontend

- React (UI library)
- Next.js 15 (full-stack framework)
- TypeScript (type-safe development)

Mobile

- Flutter (cross-platform native apps)
- FlutterFlow (rapid prototyping)

Specialized

- TensorFlow (custom ML models)
- LangSmith (agent observability and tracing)
- Microsoft Teams API (enterprise integration)
- Microsoft Graph API (enterprise ecosystem)
- .NET/C# (Revit plugin development)
- Autodesk Revit API (architectural automation)

6. CASE STUDIES AND PORTFOLIO

Case Study 1: SEOByte - Multi-Agent SEO Platform

Client: SEOByte

Website: <https://app.seobyte.ai/>

Project: 20+ specialized AI agents for end-to-end SEO automation

Completion Date: September 2025

Duration: 5 months to production

Category: AI Apps and Integration

Situation: SEO agencies spend 1 week on comprehensive keyword research, 4 hours on content briefs, and 3 days on site audits per project. SEOByte needed an AI platform to automate the entire SEO execution pipeline, from onboarding through research, clustering, to content generation, while maintaining expert-level quality.

Complication: The challenge required orchestrating 20+ specialized AI agents with complex interdependencies, building custom topical clustering algorithms for semantic understanding, integrating RAG for client-specific knowledge bases, and implementing human-in-loop quality control, all while ensuring scalable, consistent outputs across diverse domains.

What ThinkByte Built:

- 20+ specialized agents orchestrated with LangGraph (onboarding, keyword research, SERP analysis, clustering, content generation, auditing)
- Custom topical clustering algorithms for semantic keyword grouping beyond simple string matching
- RAG system for private knowledge base integration (client docs, brand guidelines, case studies)
- Editable prompt system with human-in-loop checkpoints at critical workflow stages
- FastAPI backend + Next.js dashboard with real-time analytics and workflow visualization

Results:

- Keyword research: 1 week reduced to 30 minutes (99.6% faster)
- Content briefs: 4 hours reduced to 15 minutes (94% faster)
- Site audits: 3 days reduced to 4 hours (91% faster)
- 20+ AI agents with full observability
- 90%+ workflow automation coverage
- 80% reduction in AI API costs through optimization
- Client traffic increases: 860-3,080%

Tech Stack: Python, LangGraph, LangChain, OpenAI GPT-4, Anthropic Claude, Pinecone/Weaviate (vector DB), FastAPI, React, Next.js, TypeScript, PostgreSQL, Redis, Celery, AWS

Case Study 2: NinjaHire - AI Recruitment Platform

Client: NinjaHire

Website: <https://www.ninjahire.co/>

Project: 20+ agent supervisor architecture for end-to-end recruitment automation

Completion Date: November 2025

Duration: 6 months to production

Category: AI Apps and Integration

Situation: Recruiting teams managing 500+ applicants monthly faced overwhelming operational burden: 2 full days of manual effort per hiring cycle spent on resume screening, email follow-ups, and interview scheduling. Fragmented workflows across multiple tools created inefficiency and inconsistent candidate evaluation.

Complication: This required building a supervisor-based architecture coordinating 20+ specialized agents with complex interdependencies: intelligent task routing, multi-step workflow orchestration, quality validation, persistent memory across long-running recruitment cycles (days/weeks), bidirectional ATS sync, real-time Microsoft Teams integration, and transparent editable prompts for non-technical recruiters.

What ThinkByte Built:

- LangGraph supervisor architecture intelligently routing tasks to 20+ specialized sub-agents
- Agents for: resume parsing, skill matching, screening, shortlisting, email/SMS/LinkedIn outreach, interview scheduling, feedback collection, pipeline management
- Persistent memory with vector database for context across multi-day workflows
- Event-driven ATS integration with real-time bidirectional sync
- Microsoft Teams API integration for notifications, calendar coordination, feedback collection
- Editable prompt system allowing recruiters to customize agent behavior without code

Results:

- 90% faster candidate screening (2 days reduced to 2 hours manual effort)
- 60% cost per hire reduction
- 80% automation of recruitment operations
- 10x candidate processing volume (500+ applicants/month)
- Single supervisor orchestrating 20+ specialized agents
- Seamless Microsoft Teams integration
- Editable prompts for recruiter customization
- Full workflow transparency with LangSmith tracing

Tech Stack: LangGraph, LangChain, LangSmith, OpenAI GPT-4, Pinecone/Weaviate, Python, FastAPI, Celery, Redis, PostgreSQL, React, Next.js, TypeScript, Microsoft Teams API, Microsoft Graph API

Case Study 3: Anvayaa - AI Dementia Care App (Patent-Pending)

Client: Anvayaa

Website: <http://dementiacare.anvayaa.com/>

Project: Patent-pending healthcare platform with 6+ AI systems

Completion Date: June 2025

Duration: 6 months to production

Category: AI Apps and Integration, Healthcare

Situation: 8.8 million people in India live with dementia, creating massive caregiver burden for families managing daily care, behavioral changes, and fragmented healthcare coordination. Anvayaa needed a comprehensive mobile platform (iOS + Android) to reduce caregiver stress while improving patient engagement and quality of life.

Complication: The challenge required building 6+ AI systems for personalized care: adaptive activity generation, memory box with AI photo organization and story generation, behavioral prediction models, quality of life tracking with predictive analytics, PWD profiling, and real-time care recommendations, all with elderly-accessible UX, offline functionality, multilingual support, and HIPAA-equivalent compliance.

What ThinkByte Built:

- Cross-platform Flutter app (iOS + Android) with offline-first architecture
- 6+ AI systems using LangChain/LangGraph + GPT-4 for personalized content generation
- Custom TensorFlow models for behavioral prediction and intervention recommendations
- Memory Box AI with photo analysis, story generation, and video compilation
- Real-time backend (Firebase + Supabase) with multi-device sync and caregiver coordination

- Elderly-accessible UI with voice interface and 6+ language support

Results:

- 40% reduction in caregiver stress and burden
- 5/5 star rating with positive family testimonials
- Patent application filed
- 6+ AI systems: adaptive activities, memory box, behavioral predictions, quality tracking
- 3x daily patient engagement with personalized games

Tech Stack: Flutter, Dart, LangChain, LangGraph, OpenAI GPT-4, TensorFlow, Firebase, Supabase, PostgreSQL, Google Vision API, Node.js, Python

Case Study 4: Piaxis - AI Architectural Automation (Revit Plugin)

Client: Piaxis

Website: <https://piaxis.ai/>

Project: Revit plugin with AI-powered architectural detailing

Completion Date: October 2025

Duration: 4 months to production (4-person team)

Category: AI Apps and Integration, Construction/Architecture

Situation: Architectural firms spend 30-40% of project time on repetitive construction documentation in Autodesk Revit. Piaxis needed a production-ready AI platform to automate architectural detailing while maintaining professional quality and firm-specific standards.

Complication: The challenge required real-time Revit API integration without workflow disruption, AI models trained on architectural conventions, support for multiple firm graphic standards, and enterprise-grade cloud infrastructure, all with zero data export requirements for security.

What ThinkByte Built:

- .NET/C# Revit plugin with intuitive toolbar and Design Automation API integration
- Python AI engine using LangGraph + Weaviate for intelligent detail generation
- FastAPI backend with multi-tenant architecture on AWS infrastructure
- Next.js web platform for firm administration and analytics

Results:

- 80% detailing time reduction (1 week reduced to 8 hours per project)
- 5-20 architectural firms actively using the platform
- Detail generation with real-time Revit sync
- 500-1,000+ annual hours saved per firm
- 99.5% system uptime with enterprise reliability

Tech Stack: Python, LangGraph, Weaviate, FastAPI, .NET/C#, React, Next.js, TypeScript, AWS (EC2, Lambda, S3), PostgreSQL, Redis, Docker

7. PRICING AND COMMERCIAL TERMS

Project-Based Pricing (USD)

POC/Pilot: \$5,000 to \$10,000 / 2-4 weeks / Prove concept, single feature

Small: \$10,000 to \$25,000 / 4-8 weeks / MVP, limited scope

Medium: \$25,000 to \$50,000 / 2-4 months / Full product, multiple features

Large: \$50,000 to \$100,000 / 4-6 months / Complex system, integrations

Enterprise: \$100,000+ / 6+ months / Multi-system, ongoing

Hourly Rates (USD)

Senior Engineer: \$75 to \$100 per hour (core development)

Lead/Architect: \$100 to \$150 per hour (architecture, complex problems)

PM/Coordinator: \$50 to \$75 per hour (project management)

Minimum: 10 hours for hourly engagements

Managed Teams Pricing (INR Monthly)

Backend Engineer 1 (Junior-mid): INR 1,50,000 to 1,75,000 per month

Backend Engineer 2 (Mid-senior): INR 1,75,000 to 2,25,000 per month

Backend Engineer 3 (Senior): INR 2,25,000 to 2,75,000 per month

Full Stack 1 (Junior-mid): INR 1,50,000 to 1,75,000 per month

Full Stack 2 (Mid-senior): INR 1,75,000 to 2,25,000 per month

Full Stack 3 (Senior): INR 2,25,000 to 2,75,000 per month

AI/ML Engineer (Specialized): INR 2,00,000 to 3,00,000 per month

Tech Lead (Leadership): INR 2,50,000 to 3,50,000 per month

Architect (System design): INR 3,00,000 to 4,00,000 per month

PM (Coordination): INR 1,50,000 to 2,50,000 per month

Note: +18% GST on all INR pricing

Sample Team Configurations (Monthly INR)

Starter (1 Full Stack + PM): INR 3 to 4 Lakhs per month - best for small ongoing projects

Standard (2 Engineers + PM): INR 4.5 to 5.5 Lakhs per month - best for product development

Growth (3 Engineers + Tech Lead + PM): INR 7 to 9 Lakhs per month - best for full product teams

Enterprise (4+ Engineers + Architect + PM): INR 10 Lakhs+ per month - best for large-scale systems

Payment Terms

Monthly FTE: Invoice on 25th, due Net 7, late fee 1.5% per month

Hourly Retainer: Weekly, 2 business days, work pauses if not paid

Project-Based: Milestone-based, per contract terms

Accepted payment methods: Wire transfer, ACH (US), NEFT/RTGS (India)

Not accepted: Credit cards for ongoing work, crypto, PayPal for large amounts

Discounting Guidelines

First project with strategic client: 10-15% max discount

Multi-month commitment (6+ months): 5-10% max discount

Referral from existing client: 5-10% max discount

Case study participation: 10-15% max discount

Early adopter (platform beta): 10% max discount

Pricing Red Flags (Walk Away Signals)

- Budget under \$3K for production work: Cannot deliver quality
- Wants fixed price with undefined scope: Recipe for conflict
- Expects 24/7 availability at standard rates: Unsustainable
- Compares to Fiverr/lowest-cost options: Wrong customer fit
- Won't commit to minimum engagement: Not a serious buyer

8. KEY DIFFERENTIATORS AND COMPETITIVE POSITIONING

Differentiator 1: Managed Team vs. Solo Freelancer

Problem: Clients hire freelancers who ghost, miss deadlines, or can't handle complexity.

ThinkByte Solution: Full team assigned to project, agency-level accountability, no single point of failure, consistent progress and communication.

Key Message: "Get the speed of a freelancer with the reliability of an agency."

Differentiator 2: LangGraph Specialization

Problem: Most AI developers are generalists who learn on your dime.

ThinkByte Solution: LangGraph is all they do. Deep expertise in agent architectures. Know the patterns, pitfalls, and best practices. Production-ready from day one.

Key Message: "We specialize in LangGraph - it's all we do."

Differentiator 3: Speed to Production

Problem: AI projects get stuck in "development hell" for months.

ThinkByte Solution: Proven workflows and templates. No learning curve on common patterns. Start projects in days, not weeks. Focus on unique requirements, not boilerplate.

Key Message: "Your project started in days, not weeks of hiring."

Differentiator 4: Production Experience

Problem: Many AI developers only know demos and prototypes.

ThinkByte Solution: 20+ production deployments. Handle scale, errors, edge cases. Infrastructure and DevOps experience. Monitoring and maintenance included.

Key Message: "We've taken 20+ AI systems from idea to production."

Competitive Positioning

vs. Solo Freelancers: "Unlike a solo freelancer, we're a managed team. You get multiple people working on your project, with built-in accountability."

vs. Large Agencies: "We're small and focused. No account managers or overhead - you work directly with the engineers building your system."

vs. Offshore Dev Shops: "We're US-based with overlapping hours. No communication lag or timezone issues."

vs. Internal Hiring: "Skip the 3-month hiring process. We can start this week with an experienced team."

9. VALUE PROPOSITIONS AND MESSAGING

Pain Point Messaging

For "Can't Find Talent": "Finding LangGraph developers is hard - we have a bench ready. No more interviewing 50

candidates."

For "Previous Freelancer Failed": "Managed team = accountability. We don't ghost."

For "Need to Move Fast": "Start this week, not next month. We've done this before."

For "Uncertain About Approach": "We've built 20+ LangGraph projects - we know what works and what doesn't."

For "Budget Conscious": "Managed team at freelancer rates. No agency markup."

For "Need Ongoing Support": "We don't disappear after launch. Maintenance and support included."

Credibility Statements by Project Type

RAG Systems: "We built a RAG system for an SEO platform that powers 20+ agents with client-specific knowledge bases. The key was custom topical clustering algorithms for semantic keyword grouping."

Multi-Agent Systems: "Our 20+ agent recruitment platform reduced screening from 2 days to 2 hours. The supervisor architecture intelligently routes tasks to specialized sub-agents."

Voice AI: "We've built voice pipelines with sub-500ms latency. The trick is optimizing the handoff between speech recognition and LLM processing."

Chatbots: "Our SEO platform's AI handles complex multi-step workflows with human-in-loop checkpoints. Built on LangGraph for reliability and full observability via LangSmith."

Credibility Statements by Industry

Healthcare: "We built a patent-pending dementia care app with 6+ AI systems including behavioral prediction models. We understand HIPAA considerations and elderly-accessible UX."

HR/Recruitment: "Our AI recruitment platform automates 80% of hiring operations with Microsoft Teams integration. 90% faster screening, 60% cost reduction."

Marketing/SEO: "We built SEOByte - 20+ agents that deliver 860-3,080% traffic increases. Keyword research went from 1 week to 30 minutes."

Construction/Architecture: "Our Piaxis platform reduces architectural detailing time by 80% with real-time Revit integration. 5-20 firms using it in production."

ROI Talking Points

SEOByte: 99.6% faster keyword research, 860% traffic growth

NinjaHire: 90% faster screening, 60% cost per hire reduction

Piaxis: 80% time reduction, 500-1,000 hours saved per firm per year

Anvayaa: Patent filed, 40% caregiver stress reduction

Client Praise Points

- Production-grade AI systems, not just POCs or demos
- Deep multi-agent expertise (supervisor patterns, 20+ agent orchestration)
- Healthcare innovation resulting in patent applications
- Enterprise integration mastery (Teams, Slack, ATS, Revit)
- Measurable business outcomes (90% time savings, 60% cost reduction)
- Rapid delivery (3-6 months from concept to production)
- Technical depth in both AI/ML and full-stack engineering
- Team consistency with dedicated engineers throughout project lifecycle

10. INDUSTRIES SERVED

Healthcare: Example project Anvayaa (Dementia Care). Key capabilities include HIPAA-compliant AI, patent-pending

innovation, TensorFlow behavioral models.

HR/Recruitment: Example project NinjaHire. Key capabilities include supervisor-based agent architecture, Microsoft Teams integration, ATS sync.

Marketing/SEO: Example project SEOByte. Key capabilities include multi-agent content automation, RAG knowledge bases, topical clustering.

Construction/Architecture: Example project Piaxis. Key capabilities include Revit plugin, .NET/C# integration, Design Automation API.

E-commerce: General consulting. Key capabilities include customer experience optimization, process automation.

Startups: Various projects. Key capabilities include MVP development, rapid prototyping, agent-first architectures.

Enterprise: Various projects. Key capabilities include framework selection, architecture consulting, system rescue.

11. CONTACT AND COMPANY DETAILS

Website: <https://thinkbyte.ai/>

Blog: <https://thinkbyte.ai/blog/>

LinkedIn (Company): <https://linkedin.com/company/thinkbyte.ai>

Registered Address: No. 235, 13th Cross, Hoysala Nagar 2nd Stage, Bangalore North, Bangalore, Karnataka 560038, India

Co-Founder: SriHarsha Thota - <https://www.linkedin.com/in/sriharsha-thota-37033616/>

Co-Founder: Vinay Kumar - <https://www.linkedin.com/in/vinay-kumar-2b42a845/>

Tracxn Profile: <https://tracxn.com/d/companies/thinkbyte.ai/>

Inc42 Profile: <https://inc42.com/person/sri-harsha-thota/>

12. SOURCES

- ThinkByte.AI Official Website: <https://thinkbyte.ai/>

- ThinkByte.AI Tracxn Company Profile: <https://tracxn.com/d/companies/thinkbyte.ai/>

- CX Quest Article - ThinkByte Balancing Technology and Real-World Needs: <https://cxquest.com/thinkbyte-balancing-technology-and-real-world-needs/>

- SriHarsha Thota LinkedIn: <https://www.linkedin.com/in/sriharsha-thota-37033616/>

- SriHarsha Thota Launch Post on LinkedIn: https://www.linkedin.com/posts/sriharsha-thota-37033616_thinkbyte-ai-innovation-activity-7224392390973739008-IoF6

- SriHarsha Thota Weekday.works Profile: <https://www.weekday.works/people/sriharsha-thota-sriharsha-thota-37033616>

- Sri Harsha Thota Inc42 Profile: <https://inc42.com/person/sri-harsha-thota/>

- Vinay Kumar LinkedIn: <https://www.linkedin.com/in/vinay-kumar-2b42a845/>

- ThinkByte.AI LinkedIn Company Page: <https://in.linkedin.com/company/thinkbyte.ai>

- ThinkByte Blog: <https://thinkbyte.ai/blog/>

- Internal project documentation and case study files

Document compiled: February 9, 2026

Based on publicly available information and project documentation.